

# Vehicle Insurance Prediction

```
In [59]: import pandas as pd
from hdfs import InsecureClient
import os
from pyspark.sql import SparkSession
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [4]: sparkSession = SparkSession.builder.appName("EDA Vehicle Insurance").getOrCreate()
client_hdfs = InsecureClient('hdfs://localhost:9820')
```

```
In [5]: df = sparkSession.read.csv('hdfs://localhost:9820/test/Merge.csv', header=True
, inferSchema= True)
df.show()
```

```
+---+-----+---+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+
| id|Gender|Age|Driving_License|Region_Code|Previously_Insured|Vehicle_Age|Ve
hicle_Damage|Annual_Premium|Policy_Sales_Channel|Vintage|Response|
+---+-----+---+-----+-----+-----+-----+-----+
| 1| Male| 44|          1|          28|          0| > 2 Years|
Yes|          40454|          26|          217|          1|
| 2| Male| 76|          1|          3|          0| 1-2 Year|
No|          33536|          26|          183|          0|
| 3| Male| 47|          1|          28|          0| > 2 Years|
Yes|          38294|          26|          27|          1|
| 4| Male| 21|          1|          11|          1| < 1 Year|
No|          28619|          152|          203|          0|
| 5|Female| 29|          1|          41|          1| < 1 Year|
No|          27496|          152|          39|          0|
| 6|Female| 24|          1|          33|          0| < 1 Year|
Yes|          2630|          160|          176|          0|
| 7| Male| 23|          1|          11|          0| < 1 Year|
Yes|          23367|          152|          249|          0|
| 8|Female| 56|          1|          28|          0| 1-2 Year|
Yes|          32031|          26|          72|          1|
| 9|Female| 24|          1|          3|          1| < 1 Year|
No|          27619|          152|          28|          0|
| 10|Female| 32|          1|          6|          1| < 1 Year|
No|          28771|          152|          80|          0|
| 11|Female| 47|          1|          35|          0| 1-2 Year|
Yes|          47576|          124|          46|          1|
| 12|Female| 24|          1|          50|          1| < 1 Year|
No|          48699|          152|          289|          0|
| 13|Female| 41|          1|          15|          1| 1-2 Year|
No|          31409|          14|          221|          0|
| 14| Male| 76|          1|          28|          0| 1-2 Year|
Yes|          36770|          13|          15|          0|
| 15| Male| 71|          1|          28|          1| 1-2 Year|
No|          46818|          30|          58|          0|
| 16| Male| 37|          1|          6|          0| 1-2 Year|
Yes|          2630|          156|          147|          1|
| 17|Female| 25|          1|          45|          0| < 1 Year|
Yes|          26218|          160|          256|          0|
| 18|Female| 25|          1|          35|          1| < 1 Year|
No|          46622|          152|          299|          0|
| 19| Male| 42|          1|          28|          0| 1-2 Year|
Yes|          33667|          124|          158|          0|
| 20|Female| 60|          1|          33|          0| 1-2 Year|
Yes|          32363|          124|          102|          1|
+---+-----+---+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+
only showing top 20 rows
```

```
In [7]: result_df = df.select("*").toPandas()  
result_df.head()
```

Out[7]:


	id	Gender	Age	Driving_License	Region_Code	Previously_Insured	Vehicle_Age	Vehicle_Dai
0	1	Male	44	1	28	0	> 2 Years	
1	2	Male	76	1	3	0	1-2 Year	
2	3	Male	47	1	28	0	> 2 Years	
3	4	Male	21	1	11	1	< 1 Year	
4	5	Female	29	1	41	1	< 1 Year	



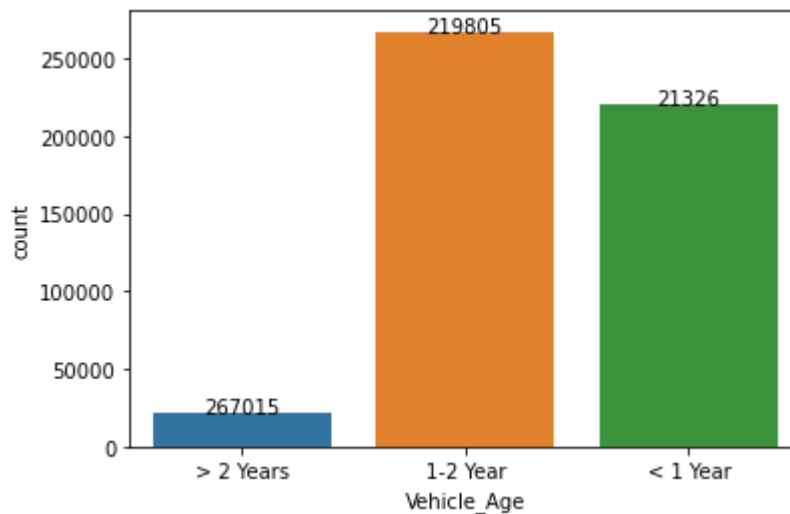
```
In [5]: print(result_df.dtypes)
print(result_df.isnull().any())
print(result_df.isnull().sum())
```

```
id                int32
Gender            object
Age              int32
Driving_License   int32
Region_Code       int32
Previously_Insured int32
Vehicle_Age       object
Vehicle_Damage    object
Annual_Premium    int32
Policy_Sales_Channel int32
Vintage           int32
Response          int32
dtype: object
id                False
Gender            False
Age              False
Driving_License   False
Region_Code       False
Previously_Insured False
Vehicle_Age       False
Vehicle_Damage    False
Annual_Premium    False
Policy_Sales_Channel False
Vintage           False
Response          False
dtype: bool
id                0
Gender            0
Age              0
Driving_License   0
Region_Code       0
Previously_Insured 0
Vehicle_Age       0
Vehicle_Damage    0
Annual_Premium    0
Policy_Sales_Channel 0
Vintage           0
Response          0
dtype: int64
```

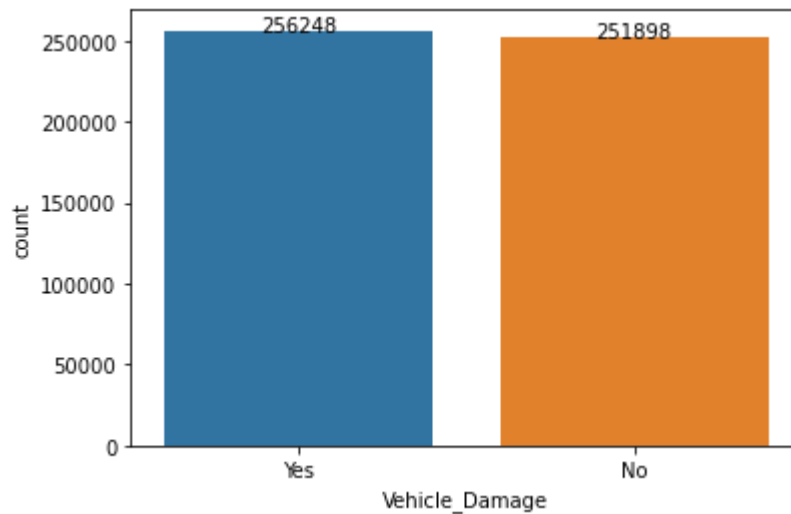
```
In [6]: graph = sns.countplot(result_df.Gender)
graph.set_xticklabels(graph.get_xticklabels())
i=0
for p in graph.patches:
    height = p.get_height()
    graph.text(p.get_x()+p.get_width()/2., height + 0.1,
               result_df['Gender'].value_counts()[i],ha="center")
    i += 1
```



```
In [7]: graph = sns.countplot(result_df.Vehicle_Age)
graph.set_xticklabels(graph.get_xticklabels())
i=0
for p in graph.patches:
    height = p.get_height()
    graph.text(p.get_x()+p.get_width()/2., height + 0.1,
               result_df['Vehicle_Age'].value_counts()[i],ha="center")
    i += 1
```




```
In [8]: graph = sns.countplot(result_df.Vehicle_Damage)
graph.set_xticklabels(graph.get_xticklabels())
i=0
for p in graph.patches:
    height = p.get_height()
    graph.text(p.get_x()+p.get_width()/2., height + 0.1,
               result_df['Vehicle_Damage'].value_counts()[i],ha="center")
    i += 1
```




```
In [9]: sns.distplot(result_df.Age)
```

```
Out[9]: <matplotlib.axes._subplots.AxesSubplot at 0x1e61a0c7730>
```



```
In [10]: sns.distplot(result_df.Annual_Premium)
```

```
Out[10]: <matplotlib.axes._subplots.AxesSubplot at 0x1e61465ce20>
```



```
In [11]: a=result_df.groupby(['Vehicle_Damage', 'Response']).size().sort_values(ascending=False).reset_index(name='Sum of Response')
b=result_df.groupby(['Previously_Insured', 'Response']).size().sort_values(ascending=False).reset_index(name='Sum of Response')
c=result_df.groupby(['Gender', 'Response']).size().sort_values(ascending=False).reset_index(name='Sum of Response')
print(a)
print(b)
print(c)
```

	Vehicle_Damage	Response	Sum of Response
0	No	0	250916
1	Yes	0	210520
2	Yes	1	45728
3	No	1	982

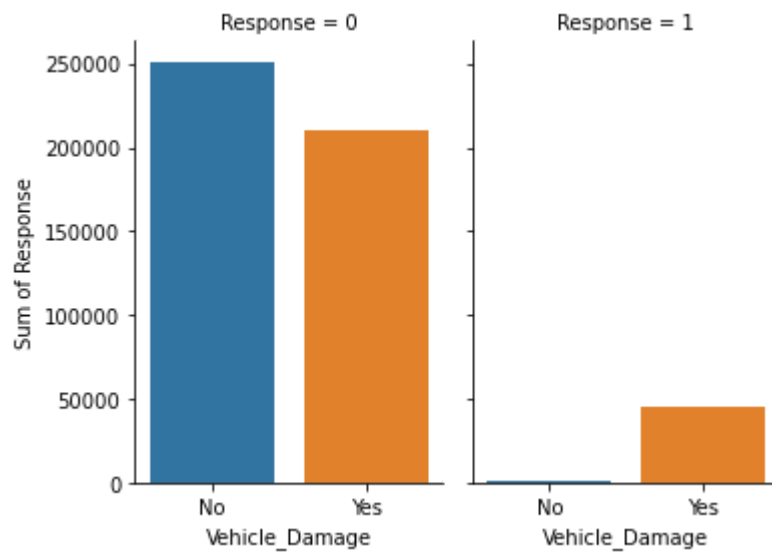
  

	Previously_Insured	Response	Sum of Response
0		1	0
1		0	0
2		0	1
3		1	1

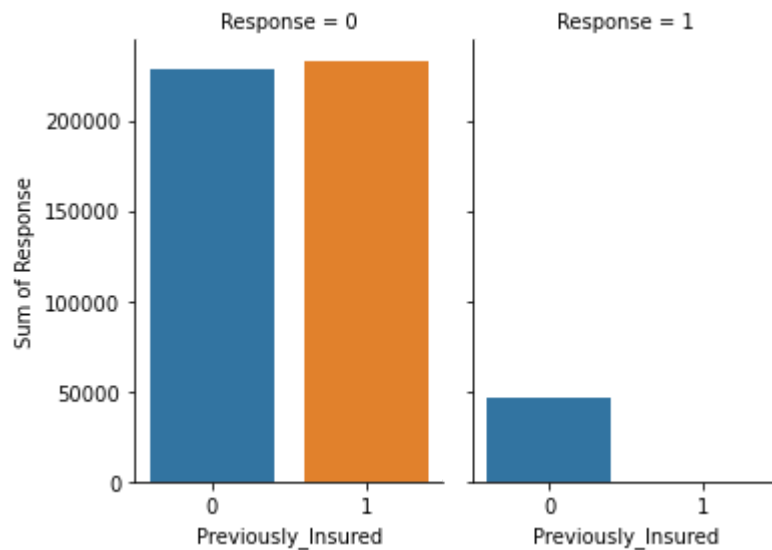
  

	Gender	Response	Sum of Response
0	Male	0	245800
1	Female	0	215636
2	Male	1	28525
3	Female	1	18185

```
In [12]: graph = sns.catplot(x="Vehicle_Damage", y="Sum of Response", col="Response",  
                             data=a, kind="bar",  
                             height=4, aspect=.7);
```




```
In [13]: graph = sns.catplot(x="Previously_Insured", y="Sum of Response", col="Response",  
                             ,  
                             data=b, kind="bar",  
                             height=4, aspect=.7);
```





```
In [14]: graph = sns.catplot(x="Gender", y="Sum of Response", col="Response",  
                             data=c, kind="bar",  
                             height=4, aspect=.7);
```



```
In [15]: result_df.Vehicle_Age.unique()
```

```
Out[15]: array(['> 2 Years', '1-2 Year', '< 1 Year'], dtype=object)
```

```
In [8]: result_df.Gender[result_df.Gender == 'Male'] = 1
result_df.Gender[result_df.Gender == 'Female'] = 0

result_df.Vehicle_Damage[result_df.Vehicle_Damage == 'Yes'] = 1
result_df.Vehicle_Damage[result_df.Vehicle_Damage == 'No'] = 0

result_df['Gender'] = result_df['Gender'].astype(int)
result_df['Vehicle_Damage'] = result_df['Vehicle_Damage'].astype(int)
```

<ipython-input-8-8f978a608a47>:1: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
result_df.Gender[result_df.Gender == 'Male'] = 1
<ipython-input-8-8f978a608a47>:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
result_df.Gender[result_df.Gender == 'Female'] = 0
<ipython-input-8-8f978a608a47>:4: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
result_df.Vehicle_Damage[result_df.Vehicle_Damage == 'Yes'] = 1
<ipython-input-8-8f978a608a47>:5: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
result_df.Vehicle_Damage[result_df.Vehicle_Damage == 'No'] = 0
```

```
In [17]: result_df.head()
```

Out[17]:

	id	Gender	Age	Driving_License	Region_Code	Previously_Insured	Vehicle_Age	Vehicle_Da
0	1	1	44	1	28	0	> 2 Years	
1	2	1	76	1	3	0	1-2 Year	
2	3	1	47	1	28	0	> 2 Years	
3	4	1	21	1	11	1	< 1 Year	
4	5	0	29	1	41	1	< 1 Year	



```
In [9]: from sklearn.preprocessing import LabelEncoder
labelencoder = LabelEncoder()
result_df['Vehicle_Age'] = labelencoder.fit_transform(result_df['Vehicle_Age'])
```

In [19]: `result_df.head()`

Out[19]:

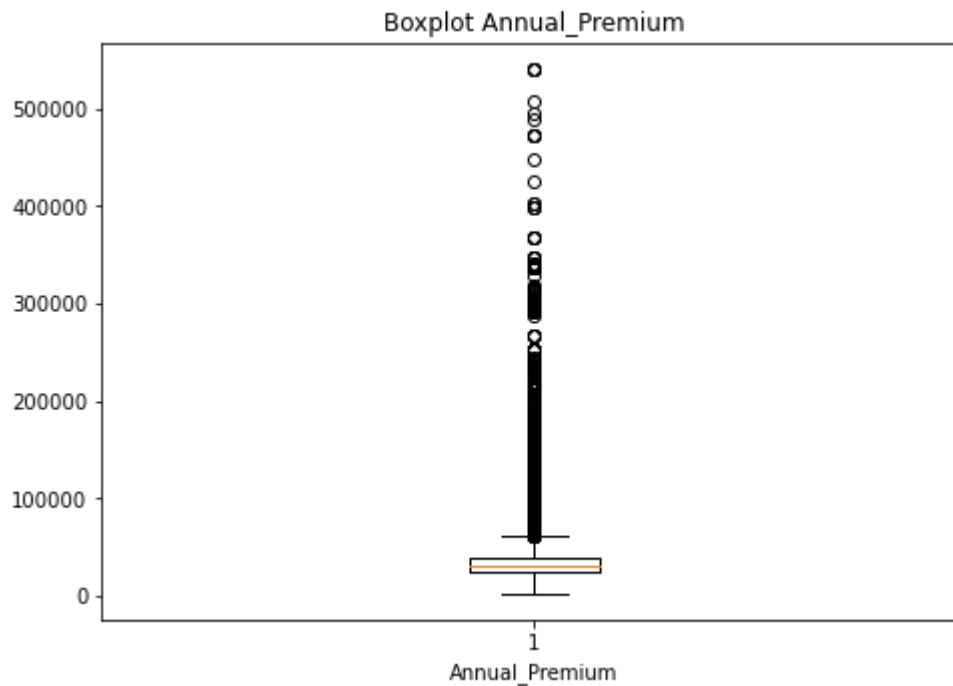
	id	Gender	Age	Driving_License	Region_Code	Previously_Insured	Vehicle_Age	Vehicle_Dai
0	1	1	44	1	28	0	2	
1	2	1	76	1	3	0	0	
2	3	1	47	1	28	0	2	
3	4	1	21	1	11	1	1	
4	5	0	29	1	41	1	1	

In [20]: `numerical_columns=['Age', 'Region_Code', 'Annual_Premium', 'Vintage']`  
`result_df[numerical_columns].describe()`

Out[20]:

	Age	Region_Code	Annual_Premium	Vintage
<b>count</b>	508146.000000	508146.000000	508146.000000	508146.000000
<b>mean</b>	38.808413	26.406572	30554.453041	154.340123
<b>std</b>	15.500179	13.224921	17146.574625	83.668793
<b>min</b>	20.000000	0.000000	2630.000000	10.000000
<b>25%</b>	25.000000	15.000000	24381.000000	82.000000
<b>50%</b>	36.000000	28.000000	31661.000000	154.000000
<b>75%</b>	49.000000	35.000000	39403.750000	227.000000
<b>max</b>	85.000000	52.000000	540165.000000	299.000000

```
In [21]: import matplotlib.pyplot as plt
fig = plt.figure()
# Create an axes instance
ax = fig.add_axes([0,0,1,1])
# Create the boxplot
bp = ax.boxplot(result_df['Annual_Premium'])
plt.xlabel("Annual_Premium")
plt.title("Boxplot Annual_Premium")
plt.show()
```



```
In [22]: import numpy as np
outliers=[]
def detect_outlier(data_1):

    threshold=3
    mean_1 = np.mean(data_1)
    std_1 =np.std(data_1)

    for y in data_1:
        z_score= (y - mean_1)/std_1
        if np.abs(z_score) > threshold:
            outliers.append(y)
    return outliers
```

```
In [23]: outlier_datapoints = detect_outlier(result_df["Annual_Premium"])  
print(outlier_datapoints)
```

[89282, 101021, 82705, 90789, 119735, 104002, 92716, 112974, 139130, 98002, 267698, 125643, 85786, 95217, 101064, 86283, 136061, 117799, 84142, 91520, 87273, 133098, 103026, 87831, 152331, 100688, 101069, 87954, 131469, 104529, 86336, 82081, 90024, 98428, 107966, 89687, 137771, 82470, 88368, 88577, 90967, 91670, 83846, 99793, 101048, 90046, 95554, 88825, 82932, 123745, 103372, 84943, 103758, 101716, 508073, 89480, 88129, 101904, 94900, 90526, 83368, 82085, 93104, 84472, 82144, 141770, 95569, 301762, 120037, 119148, 84077, 95258, 110204, 89463, 315565, 93104, 85772, 89902, 85670, 126671, 107748, 92211, 87105, 86793, 83912, 107266, 84056, 89637, 124345, 98425, 82879, 91565, 113820, 89355, 104781, 91407, 95895, 89888, 183718, 98526, 88378, 86983, 89884, 86606, 106578, 116045, 86416, 95598, 94647, 82422, 110973, 99324, 94506, 93656, 87578, 95221, 109361, 86885, 92190, 95817, 84122, 83433, 103906, 87763, 98337, 82876, 181076, 111257, 160380, 159869, 229935, 94109, 88281, 83137, 86646, 99999, 84475, 90305, 94580, 140448, 113810, 82141, 82231, 82067, 82420, 87128, 98979, 92575, 99851, 87307, 89687, 147075, 98152, 94333, 98073, 83196, 84432, 101568, 100018, 168597, 82135, 119425, 97229, 83778, 101024, 85327, 83666, 86711, 91713, 155317, 83045, 90374, 99529, 160011, 151585, 113332, 199996, 96639, 111070, 104511, 113339, 105335, 83951, 294209, 83017, 104998, 86750, 82125, 94216, 90834, 107454, 85608, 86415, 89440, 336395, 95824, 83037, 82753, 89778, 95992, 111447, 95641, 214455, 82701, 89388, 95692, 88883, 87449, 96374, 93597, 86349, 86484, 97164, 90040, 86897, 85198, 89721, 111327, 101935, 88891, 88746, 84472, 88038, 97467, 109546, 94843, 82128, 253362, 141131, 85340, 99873, 83203, 82424, 91146, 117793, 135054, 82089, 98327, 82033, 98152, 92928, 97936, 113379, 102708, 102724, 88196, 99096, 84917, 116398, 89263, 88123, 95664, 99198, 93267, 84918, 91694, 90534, 211132, 174574, 89022, 83652, 89425, 101123, 296891, 106843, 83957, 90972, 117563, 97505, 83338, 128329, 113088, 107115, 91151, 122700, 189361, 235683, 98067, 105627, 87665, 92716, 101331, 540165, 104003, 92518, 103511, 169127, 98650, 97693, 121384, 89063, 91415, 85063, 88638, 92860, 103968, 82129, 83808, 85553, 127493, 86860, 99950, 101025, 100196, 86958, 84875, 95217, 93259, 340439, 82501, 105542, 97772, 123446, 87557, 95904, 99238, 101825, 89388, 100993, 229375, 83638, 88673, 113636, 84328, 129892, 88134, 95692, 91118, 102962, 85000, 121448, 82372, 252141, 90834, 142313, 206906, 106505, 131355, 84268, 85207, 83958, 108034, 91338, 107971, 84025, 218552, 82491, 86205, 240318, 105714, 101333, 83228, 111486, 135616, 110070, 98905, 85480, 86426, 86581, 82237, 167393, 83876, 120406, 126410, 107103, 107943, 99091, 84810, 92161, 167899, 142271, 113029, 152866, 91323, 93214, 109738, 94946, 84671, 291169, 104539, 150468, 143525, 86013, 99248, 91164, 117237, 100718, 133747, 83683, 126158, 102747, 113553, 85514, 101380, 104388, 94899, 240809, 103439, 88128, 99339, 91487, 83182, 103059, 101639, 93971, 109080, 214455, 104195, 182288, 85990, 94343, 171264, 96436, 88946, 88533, 92454, 241735, 111384, 107115, 188591, 318706, 114021, 82195, 85683, 93902, 84046, 95274, 96364, 88388, 82649, 84112, 124520, 82204, 88369, 267698, 93195, 98418, 131901, 99647, 87798, 86112, 164502, 131148, 90959, 106829, 91882, 103417, 83076, 93569, 336395, 88908, 83646, 127942, 88411, 91886, 87601, 83817, 83249, 82968, 89059, 84025, 100985, 251853, 94593, 100377, 110308, 82788, 118594, 83554, 87133, 110498, 199154, 92267, 135229, 92300, 82826, 106909, 87675, 87130, 101771, 85787, 110410, 92064, 92976, 112634, 99445, 84456, 313854, 109161, 110686, 82613, 84008, 88518, 94770, 88690, 82753, 102867, 84945, 88183, 121760, 165501, 94391, 93902, 88664, 134896, 286666, 88434, 98794, 159995, 111419, 83249, 86047, 95840, 90446, 90940, 87263, 94843, 97533, 83826, 309867, 84008, 84514, 84046, 83027, 90600, 129706, 117544, 88811, 100278, 90335, 94035, 86343, 104229, 91094, 99866, 106392, 336395, 122795, 83979, 89643, 87386, 192034, 87864, 84887, 84901, 93931, 199154, 101661, 303550, 100894, 86719, 87921, 114873, 99764, 118594, 102815, 82972, 108765, 100691, 93804, 90526, 84018, 97966, 102552, 95411, 96457, 85949, 125475, 87988, 82115, 102917, 95021, 89355, 97514, 112816, 128993, 86776, 120595, 211132, 313424, 84725, 83027, 103723, 113350, 89408, 82202, 108305, 99171, 86968, 92717, 139556, 217801, 123641, 84077,

94435, 92298, 93950, 92396, 101045, 87702, 140912, 116796, 84918, 89857, 84582, 83086, 199303, 106670, 84609, 188591, 98893, 313854, 101070, 96120, 85851, 86349, 86659, 105998, 105637, 86624, 93335, 108182, 89237, 91625, 88522, 84399, 90235, 99449, 130427, 97733, 156270, 83096, 92268, 82354, 139893, 87760, 448156, 83817, 156062, 82272, 87575, 88366, 97865, 86631, 85525, 84223, 107566, 100549, 93664, 104884, 83771, 102975, 82519, 112113, 87651, 88940, 109461, 96344, 91599, 97823, 104078, 99528, 104643, 184270, 86987, 94770, 183718, 89237, 86592, 85760, 141582, 91407, 114393, 112667, 104726, 135653, 336395, 83846, 109046, 103890, 99887, 105538, 89311, 83869, 152895, 313424, 96532, 214330, 308615, 104643, 128686, 97803, 123233, 295106, 88995, 84715, 88446, 85209, 109865, 173513, 82968, 95165, 84046, 92440, 87403, 121207, 82195, 87556, 91531, 472042, 90306, 82569, 93629, 90515, 111779, 140293, 101888, 82753, 100254, 101338, 93427, 143007, 97293, 87006, 115719, 97518, 91459, 93012, 85991, 83102, 101382, 94810, 158307, 85040, 122625, 133280, 84025, 98417, 111683, 122917, 90736, 113177, 96713, 123247, 92584, 91897, 399010, 134013, 84025, 94516, 115034, 144209, 100806, 102789, 96135, 82613, 101215, 82657, 85196, 170758, 97223, 102997, 540165, 97868, 88072, 113350, 346200, 93995, 112785, 104233, 93259, 98529, 96074, 227281, 94892, 101802, 114452, 82802, 83013, 90544, 85235, 93691, 90861, 98662, 109279, 84196, 88523, 95895, 82162, 140225, 85517, 85751, 104153, 85706, 87160, 86887, 96997, 90761, 91274, 111208, 174592, 85862, 126302, 90025, 88023, 103510, 111715, 103123, 104290, 83525, 313854, 89090, 89408, 115554, 93095, 105939, 102326, 82096, 89657, 134092, 88964, 100253, 85431, 126576, 84080, 101981, 101322, 90830, 86378, 93050, 93367, 89126, 85888, 97370, 195863, 98301, 88419, 299811, 109019, 83910, 91644, 88733, 107531, 84028, 90024, 144004, 86482, 83859, 85947, 99328, 97986, 84405, 106773, 90151, 114210, 152654, 124981, 97363, 94886, 90075, 82482, 93345, 118971, 124342, 103934, 85104, 83151, 117057, 89336, 115719, 109099, 97259, 98112, 83261, 90109, 84025, 97265, 90359, 107217, 84085, 89163, 91452, 111883, 82358, 89960, 83731, 99584, 82342, 105335, 83005, 106035, 251853, 84459, 82822, 122325, 91266, 336395, 87880, 115719, 84943, 107119, 87954, 85991, 112746, 93597, 316563, 85720, 100296, 84901, 97803, 162785, 88955, 94435, 96266, 85486, 89478, 83666, 489663, 86144, 107875, 96136, 293038, 126576, 105379, 115719, 91545, 92465, 173545, 93253, 83873, 83378, 234236, 83907, 96253, 92203, 87864, 118167, 82708, 86983, 97335, 100591, 114757, 111876, 95159, 199996, 87763, 97710, 92353, 82222, 121058, 84728, 107561, 113857, 97370, 110479, 95843, 82928, 86455, 84397, 101205, 87930, 88669, 86485, 87651, 92249, 101146, 87891, 100574, 83586, 103059, 89604, 97799, 170661, 127279, 89287, 84630, 83757, 90834, 112988, 88811, 91037, 110592, 96079, 93735, 110999, 101380, 95118, 88683, 121321, 86492, 91369, 83148, 87813, 83151, 83813, 101108, 86656, 84771, 85804, 86983, 82647, 98307, 100950, 143007, 104229, 110308, 83569, 90718, 366891, 84921, 98237, 82737, 110640, 96500, 86697, 91151, 92764, 540165, 109321, 116777, 106035, 88634, 85555, 87796, 85290, 91969, 91379, 94080, 101892, 133280, 88841, 132387, 96708, 96609, 99695, 89345, 82753, 346982, 119446, 85019, 120507, 84241, 93597, 120245, 92631, 86816, 90771, 88469, 107078, 158356, 85549, 110012, 91369, 85479, 84928, 87386, 173513, 83423, 94606, 95117, 109279, 101333, 91545, 83242, 104535, 93786, 82657, 87750, 114391, 100227, 93376, 91698, 102261, 102109, 82685, 87470, 100817, 177150, 107287, 82148, 111671, 86106, 101680, 83638, 105148, 94216, 88557, 93705, 86806, 87901, 100227, 83187, 83527, 135545, 113486, 214595, 107966, 96371, 96030, 88148, 109321, 91545, 98602, 106843, 87914, 82753, 96030, 100064, 85787, 84392, 97038, 97718, 87582, 89449, 94779, 93597, 122442, 106836, 95014, 110276, 92954, 103603, 101802, 366891, 96941, 133355, 83854, 144074, 85758, 83670, 97468, 89956, 118167, 100245, 90123, 162948, 110592, 93485, 196513, 92994, 127772, 142995, 86784, 84600, 95130, 229935, 100502, 87271, 84742, 110012, 111518, 329525, 91348, 233713, 83666, 83787, 109890, 87956, 91414, 93214, 95845, 86920, 89929, 89088, 117799, 89178, 82713, 93449, 90216, 347606, 85535, 108323, 95712, 87160, 129609, 185009, 86380, 105365, 102637, 91214, 87666, 108651, 118745, 107003, 88656, 90988, 91677, 129749, 1063

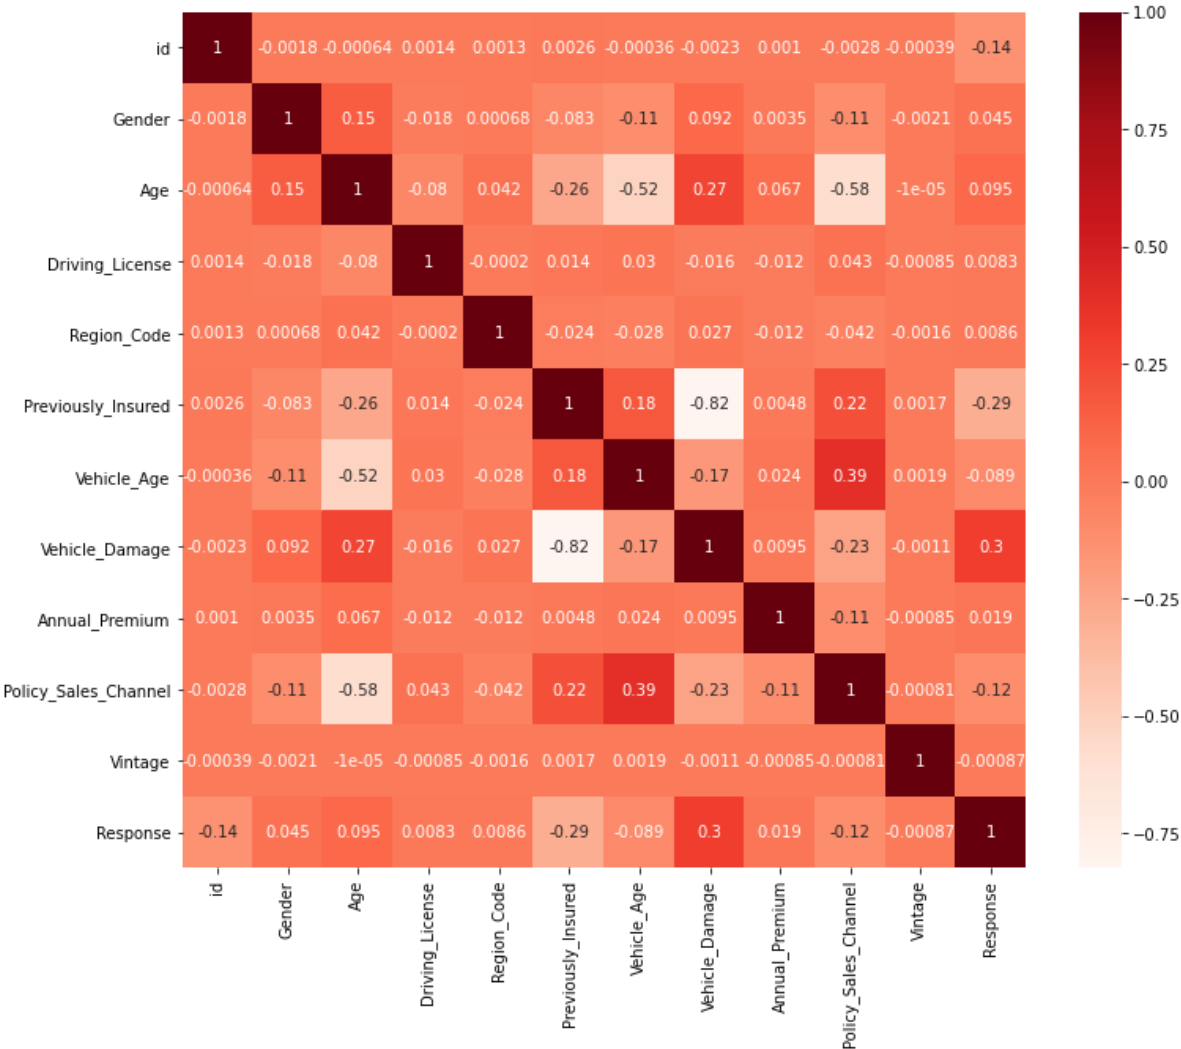
35, 90090, 91670, 93611, 83027, 95817, 87364, 90856, 84915, 82031, 84176, 954  
07, 97917, 89464, 83590, 130703, 96095, 144473, 83196, 83526, 84641, 86618, 8  
6482, 127575, 85055, 87363, 94135, 93345, 93950, 132308, 117793, 83590, 9369  
1, 96724, 86259, 82419, 85872, 95753, 113800, 93951, 99715, 101401, 109570, 1  
01190, 109292, 132387, 99445, 82342, 89139, 97979, 97325, 89334, 152009, 9189  
7, 87295, 92906, 87065, 88985, 85600, 87470, 86013, 82024, 143814, 87692, 857  
97, 86595, 89419, 103807, 96127, 86082, 93136, 117793, 113684, 101401, 88944,  
99445, 164990, 89843, 91916, 85949, 113512, 119025, 84394, 232791, 95578, 103  
415, 87386, 95875, 117237, 83646, 102757, 95135, 111040, 139130, 105842, 8263  
4, 114391, 93149, 101413, 83808, 91505, 82068, 82216, 90162, 89509, 97523, 96  
248, 90523, 103262, 88026, 152870, 402863, 98645, 97997, 90985, 89605, 98444,  
94696, 83566, 82201, 92129, 86807, 84901, 91638, 82424, 126583, 299410, 9522  
6, 98893, 117993, 89373, 82031, 102797, 91279, 110204, 86860, 94477, 96108, 9  
5971, 103828, 98073, 113684, 82657, 92291, 291169, 94826, 84380, 106671, 3363  
95, 113452, 105538, 105842, 106897, 87101, 82533, 218095, 90941, 88833, 8758  
3, 105836, 88996, 87750, 110496, 252866, 84380, 86396, 92738, 98854, 294467,  
111114, 101680, 103397, 98337, 84036, 138002, 366891, 86941, 97335, 87474, 84  
984, 82998, 88433, 83963, 107748, 87274, 107195, 131469, 110266, 84772, 9102  
3, 100772, 96067, 87084, 94942, 90006, 89377, 102045, 103068, 114628, 82176,  
89408, 95284, 96603, 216660, 119272, 91260, 110052, 95432, 98934, 88850, 8467  
0, 86320, 130272, 82409, 108768, 97413, 185009, 83666, 89101, 85568, 92790, 1  
40912, 88653, 118594, 100817, 92232, 101575, 117793, 103417, 108961, 87031, 8  
2284, 85657, 296891, 97523, 92166, 109448, 97305, 86182, 101670, 93089, 12845  
7, 82414, 86823, 101682, 82713, 89041, 90156, 117793, 174351, 82454, 83174, 8  
2753, 82295, 87228, 82375, 88018, 91606, 82682, 88363, 143445, 118613, 87545,  
84786, 99474, 99685, 101401, 100445, 105212, 89549, 101397, 83666, 96347, 905  
41, 85604, 399010, 86082, 87004, 201587, 115232, 94331, 111557, 83997, 84380,  
124950, 87120, 229935, 100337, 88883, 92978, 92994, 89067, 106843, 162733, 54  
0165, 86199, 84807, 98184, 102285, 102046, 96266, 104635, 82824, 109080, 1229  
06, 99718, 96158, 91770, 83355, 85571, 336395, 140789, 87431, 243141, 94446,  
82675, 91029, 83005, 291169, 89581, 91396, 82314, 99021, 472042, 82932, 11041  
0, 91452, 94955, 90809, 110204, 91108, 84651, 85599, 227526, 114210, 90985, 8  
5466, 94970, 82682, 114145, 92091, 91279, 84950, 103222, 88040, 104521, 9109  
4, 84109, 102314, 472042, 85049, 109236, 86045, 93389, 83124, 103748, 121601,  
92453, 93089, 93983, 87759, 88773, 83265, 99528, 134022, 91698, 85332, 87057,  
99750, 296891, 136156, 83424, 131874, 116796, 84699, 107189, 96061, 119425, 8  
2738, 91969, 110479, 92042, 114391, 97056, 119272, 96199, 118644, 108034, 828  
52, 91505, 85055, 291169, 86394, 97865, 97187, 112176, 84649, 89765, 82058, 3  
40439, 105379, 164759, 87507, 86182, 89375, 106396, 85442, 88623, 87679, 1736  
53, 84811, 86205, 84298, 85347, 86275, 108692, 91533, 116606, 86349, 155317,  
93458, 92386, 92877, 97155, 84178, 85983, 235683, 83666, 85752, 158860, 9395  
0, 96425, 92795, 84396, 85922, 91713, 84715, 99987, 107446, 115014, 311948, 8  
9422, 86084, 86943, 85268, 135998, 102370, 183718, 120134, 96482, 82877, 9044  
6, 134092, 100544, 137041, 84818, 90007, 93036, 95362, 96884, 88518, 95014, 9  
2181, 110655, 94253, 108034, 91703, 155184, 90514, 136113, 91605, 95121, 9154  
5, 82958, 112115, 94181, 88433, 107566, 109965, 89575, 94931, 97748, 99730, 8  
9090, 135450, 232791, 97151, 88806, 107120, 211132, 92555, 85787, 84007, 9262  
1, 84055, 174094, 83117, 86092, 85993, 101852, 235683, 128025, 86512, 88955,  
83654, 97391, 97510, 83979, 116796, 83196, 135038, 83233, 84807, 112024, 8963  
7, 116272, 86419, 119373, 114019, 88403, 84452, 84398, 86452, 91279, 83700, 8  
4084, 84098, 85087, 89637, 96026, 87971, 86781, 86378, 85370, 96879, 97545, 9  
2300, 92120, 88419, 108766, 86888, 105608, 82649, 339396, 91807, 88072, 10423  
7, 123396, 85722, 84944, 101013, 88124, 84824, 113663, 84085, 104710, 169021,  
98645, 83360, 93597, 104939, 84912, 82060, 109803, 82705, 118922, 89419, 1130  
88, 97065, 122051, 125876, 366891, 96230, 122051, 82928, 102046, 150412, 9146  
9, 84328, 110049, 132193, 93337, 82632, 86574, 313854, 115554, 93042, 101568,  
100543, 123783, 90559, 87498, 89646, 108768, 97070, 84844, 97869, 84004, 9324



2, 106213, 85055, 84866, 85442, 85582, 84028, 402097, 187756, 114983, 108653, 86000, 101070, 84386, 91882, 86956, 87758, 89652, 92994, 94798, 86834, 103314, 155413, 82531, 82000, 88989, 83489, 106646, 85787, 93925, 90355, 85858, 93104, 122625, 85315, 105503, 102948, 95739, 109080, 83424, 87536, 92898, 87130, 120340, 87286, 85447, 199996, 99611, 110592, 116852, 88591, 109080, 82914, 115672, 85108, 337327, 82662, 97038, 266118, 98442, 156389, 95221, 96468, 85074, 91407, 159243, 91293, 211132, 122424, 114536, 96060, 141863, 136163, 144224, 85828, 188591, 132401, 92161, 100656, 131570, 87980, 103561, 91876, 84025, 97615, 104962, 111671, 90585, 89769, 86592, 90138, 113230, 87658, 92860, 125490, 100255, 94581, 90019, 92041, 88518, 170402, 86013, 110573, 102049, 84219, 83187, 117683, 94846, 87736, 86608, 98337, 82074, 99552, 145053, 88940, 85259, 82880, 83871, 98922, 85514, 88664, 89168, 91267, 85977, 86272, 128993, 100172, 89164, 119963, 88659, 82780, 84877, 94957, 96230, 127161, 82572, 121489, 127454, 495106, 88653, 159243, 111883, 109463, 84080, 95793, 90682, 83458, 92516, 98299, 88800, 92898, 91004, 98794, 94327, 101808, 83586, 128941, 102744, 86408, 86562, 103534, 87056, 90175, 82017, 87759, 85777, 104965, 98832, 92210, 112611, 88057, 85402, 313424, 88825, 87674, 147075, 97068, 89139, 91061, 111683, 84585, 95850, 92175, 235683, 82067, 102257, 83600, 82067, 96801, 92347, 107292, 185338, 100094, 92985, 94249, 85467, 88891, 88296, 104280, 85283, 98571, 107680, 83572, 108736, 100825, 84230, 84691, 220702, 82466, 90926, 102326, 93951, 82935, 109762, 94431, 102974, 92428, 96451, 82326, 118327, 92327, 82060, 83900, 96349, 244589, 103314, 83666, 93360, 98428, 87043, 88328, 104566, 83290, 111384, 85230, 86476, 199996, 97741, 115092, 83321, 82979, 88977, 84025, 104765, 91214, 84331, 92381, 86774, 291169, 140225, 82307, 98956, 84469, 86855, 86573, 82613, 167548, 82373, 102753, 102874, 89601, 82024, 188014, 99096, 107971, 109015, 104388, 90947, 83575, 295106, 106035, 90985, 87815, 97432, 346982, 109374, 91698, 87044, 167548, 109080, 111268, 92597, 100171, 96944, 151034, 85313, 235683, 220581, 93312, 87664, 111257, 337573, 89956, 93597, 86366, 99986, 92185, 119593, 85516, 94767, 91267, 101813, 94819, 83643, 90989, 83396, 126576, 251817, 93458, 85555, 84312, 90637, 85034, 84319, 99997, 89054, 85757, 124881, 97693, 86164, 89582, 87130, 165395, 85060, 94121, 93221, 92283, 95453, 91305, 101664, 110611, 93515, 90262, 84397, 86496, 92601, 111358, 123928, 82005, 86527, 83278, 87956, 106114, 105013, 101160, 90648, 106344, 87401, 87067, 196587, 143525, 90090, 144874, 88316, 106035, 122351, 88659, 99445, 102747, 95102, 82824, 120320, 91164, 87880, 89695, 93259, 100950, 192052, 87709, 90800, 112667, 109299, 339396, 107321, 94581, 88278, 110323, 87423, 174721, 94717, 82758, 103195, 87217, 104924, 92555, 101938, 97164, 84046, 99129, 89329, 83758, 83915, 130474, 130980, 92014, 85108, 120927, 117799, 86104, 97718, 82421, 84529, 94128, 137796, 87228, 84887, 82028, 130307, 175886, 122625, 99258, 82962, 86587, 110479, 83666, 110592, 90765, 297040, 125848, 86191, 83924, 90584, 88117, 96780, 92932, 100551, 199769, 87439, 87730, 91798, 85068, 196890, 123362, 84490, 83234, 93360, 116451, 98102, 105052, 107414, 83076, 86069, 108764, 142271, 118964, 96856, 114164, 97694, 83597, 98317, 87891, 157479, 84742, 109958, 90299, 82968, 86467, 91389, 140744, 83586, 97823, 340439, 90507, 154624, 88870, 86793, 135109, 169300, 113656, 99024, 82713, 93950, 114762, 108182, 84766, 86496, 84213, 82454, 136253, 87056, 83177, 91668, 96571, 88989, 96316, 108800, 90838, 90861, 104966, 85780, 101953, 105979, 90700, 121275, 104819, 90918, 83095, 93360, 101535, 82051, 95342, 102931, 107218, 93521, 96495, 112264, 86218, 99445, 98519, 92228, 130047, 241735, 109538, 122134, 122625, 108323, 122494, 85610, 88269, 124488, 199996, 101449, 93473, 99340, 95187, 91882, 92757, 87032, 97847, 103931, 108707, 91545, 84396, 102135, 303550, 91637, 91279, 83096, 87755, 88529, 108768, 113350, 188591, 89650, 132789, 84521, 103999, 154665, 174859, 87617, 91038, 84984, 85333, 94322, 94767, 150749, 84339, 88326, 86861, 83688, 84432, 103439, 182288, 126859, 83206, 97216, 119191, 83666, 159979, 115702, 141394, 93837, 188591, 116777, 101172, 95310, 84611, 109515, 90791, 183718, 85582, 115189, 106035, 101022, 85983, 170763, 96371, 97928, 102513, 95109, 99339, 104108, 138836, 110592, 100981, 108

208, 107782, 88618, 101450, 119223, 129901, 88060, 148931, 101026, 84541, 83589, 83125, 86542, 152417, 174814, 96074, 98783, 109577, 84494, 103150, 116777, 113577, 87295, 90837, 237697, 87178, 82578, 85127, 109003, 114589, 165395, 87988, 98507, 82833, 109515, 149375, 90138, 93983, 82002, 83512, 211132, 200998, 84841, 84679, 90607, 130760, 91138, 108058, 93597, 90116, 90601, 220702, 97949, 121123, 84611, 111025, 92985, 96606, 399010, 204681, 95773, 126809, 87708, 88567, 90689, 100431, 101049, 96392, 106112, 99604, 147943, 114090, 82577, 102572, 150508, 96987, 87514, 90809, 336395, 84609, 106533, 84157, 90584, 108400, 424578, 103811, 101997, 101371, 82619, 105969, 91266, 87470, 92783, 8274, 94111, 86362, 88282, 90630, 93597, 82925, 98849, 92859, 188591, 90070, 84109, 102931, 87553, 85549, 145072, 101942, 109372, 88618, 105887, 86253, 90967, 84396, 96620, 84080, 174298, 90452, 86579, 96045, 83499, 93389, 92042, 90616, 96616, 92700, 85484, 87754, 92691, 146390, 90467, 92995, 106035, 215905, 97710, 116363, 103191, 96759, 118916, 89535, 97852, 124626, 92181, 103886, 83008, 472042, 88128, 143163, 123008, 98800, 84414, 218552, 339396, 120319, 195863, 90250, 101684, 102328, 95763, 84585, 99611, 94335, 90353, 91324, 91060, 193637, 178137, 99445, 105969, 84471, 88409, 83683, 145414, 87395, 92228, 90690, 94557, 105379, 101953, 100690, 96135, 95136, 83132, 98794, 89440, 99452, 84298, 172150, 96821, 127292, 117683, 96068, 132789, 84187, 91312, 127242, 101795, 82397, 111040, 113636, 107412, 117794, 115252, 88985, 99506, 106737, 232791, 87548, 85159, 111327, 110068, 82581, 95221, 82868, 103534, 83638, 98152, 87914, 84786, 83278, 89419, 100218, 95120, 82908, 87844, 96636, 100218, 95661, 88683, 99706, 89282, 105538, 87914, 101078, 90727, 82491, 113260, 82775, 92764, 89379, 82598, 267909, 107412, 84085, 99596, 84501, 92865, 92901, 85956, 103150, 94320, 94506, 89162, 103629, 91672, 106274, 88405, 111484, 154354, 86916, 105779, 104634, 93902, 95773, 85797, 107729, 85304, 86465, 98425, 98018, 98645, 188591, 86455, 106112, 87244, 118626, 82598, 126576, 91267, 101146, 83206, 102832, 90967, 83739, 106570, 82353, 92428, 91505, 106112, 148945, 110012, 93738, 98331, 89637, 141411, 92404, 105006, 123425, 126576, 106486, 95962, 93195, 84560, 89705, 83666, 87336, 90541, 93597, 82740, 83666, 110323, 104774, 91255, 109582, 109067, 82028, 96495, 86659, 87096, 99941, 100189, 315565, 100796, 83523, 87657, 84702, 106219, 83338, 82914, 133633, 91469, 93138, 89456, 92005, 89305, 116937, 83338, 93253, 84108, 82317, 99753, 103415, 90198, 95150, 93922, 89950, 98741, 84552, 303339, 105802, 102626, 122186, 88555, 90918, 93367, 89242, 97955, 91897, 82254, 92120, 108008, 106973, 93253, 92120, 122784, 109321, 87386, 94743, 138002, 84026, 151034, 83958, 87271, 86013, 83027, 134747, 199154, 109019, 107586, 119414, 124131, 105335, 83666, 87741, 82335, 154802, 90630, 162785, 108957, 91531, 100004, 98177, 84613, 91882, 83757, 93030, 102744, 89843, 99512, 86718, 163442, 91939, 118312, 122625, 110623, 93021, 88891, 83172, 89718, 94506, 98289, 91969, 98164, 85148, 94781, 227281]

```
In [24]: #Using Pearson Correlation
plt.figure(figsize=(12,10))
cor = result_df.corr()
sns.heatmap(cor, annot=True, cmap=plt.cm.Reds)
plt.show()
#Correlation with output variable
cor_target = abs(cor["Annual_Premium"])
#Selecting highly correlated features
relevant_features = cor_target[cor_target>0]
print("Correlation Coefficeint With Respect to Response")
print(relevant_features)
```



Correlation Coefficeint With Respect to Response

id	0.001042
Gender	0.003502
Age	0.067392
Driving_License	0.012084
Region_Code	0.012012
Previously_Insured	0.004776
Vehicle_Age	0.023545
Vehicle_Damage	0.009504
Annual_Premium	1.000000
Policy_Sales_Channel	0.114035
Vintage	0.000855
Response	0.019477

Name: Annual\_Premium, dtype: float64

# Linear Regression

```
In [25]: x=result_df.drop(['Annual_Premium','id'], axis = 1)
y=result_df['Annual_Premium']
print(x.shape)
print(y.shape)
```

```
(508146, 10)
(508146,)
```

```
In [26]: from sklearn.model_selection import train_test_split
train_features, test_features, train_labels, test_labels = train_test_split(x,
y, test_size = 0.3, random_state = 0)
```

```
In [27]: print('Training Features Shape:', train_features.shape)
print('Training Labels Shape:', train_labels.shape)
print('Testing Features Shape:', test_features.shape)
print('Testing Labels Shape:', test_labels.shape)
```

```
Training Features Shape: (355702, 10)
Training Labels Shape: (355702,)
Testing Features Shape: (152444, 10)
Testing Labels Shape: (152444,)
```

```
In [28]: from sklearn.linear_model import LinearRegression
model = LinearRegression()
model.fit(train_features, train_labels)
y_pred = model.predict(test_features)
```

```
In [29]: print("Training set score: {:.4f}".format(model.score(train_features,train_labels)))

print("Test set score: {:.4f}".format(model.score(test_features,test_labels)))
```

```
Training set score: 0.0216
Test set score: 0.0202
```

```
In [30]: print(model.coef_)
print(model.intercept_)
```

```
[-2.20296565e+02  5.68341300e+01 -2.40386630e+03 -2.20994504e+01
 1.84695901e+03  2.87446966e+03  7.89872640e+02 -4.05438669e+01
 -2.30672795e-01  9.38180788e+02]
33231.52629203656
```

```
In [31]: from sklearn.metrics import mean_squared_error
mse = mean_squared_error(test_labels, y_pred)
rmse = np.sqrt(mse)
print("RMSE value: {:.4f}".format(rmse))
```


```
RMSE value: 16885.9328
```

```
In [32]: # Calculate and print r2_score


from sklearn.metrics import r2_score
print ("R2 Score value: {:.4f}".format(r2_score(test_labels, y_pred)))
```

R2 Score value: 0.0202

```
In [33]: # Plotting residual errors
plt.scatter(model.predict(train_features), model.predict(train_features) - tra
in_labels, color = 'red', label = 'Train data')
plt.scatter(model.predict(test_features), model.predict(test_features) - test_
labels, color = 'blue', label = 'Test data')
plt.hlines(xmin = 20000, xmax = 50000, y = 0, linewidth = 3)
plt.title('Residual errors')
plt.legend(loc = 4)
plt.show()
```




```
In [34]: import statsmodels.api as sm
import pylab
import scipy.stats as stats
stats.probplot(y_pred.reshape(-1), dist="norm", plot=pylab)
pylab.show()
```



```
In [35]: import seaborn as sns
residuals = test_labels - y_pred
ax = sns.regplot(x=residuals, y=y_pred)
ax.set_title('Residual vs Fitted')
```

```
Out[35]: Text(0.5, 1.0, 'Residual vs Fitted')
```



## Lasso, Ridge

```
In [36]: from sklearn.linear_model import Lasso
```

```

In [37]: lasso = Lasso()
lasso.fit(train_features,train_labels)
train_score=lasso.score(train_features,train_labels)
test_score=lasso.score(test_features,test_labels)
coeff_used = np.sum(lasso.coef_!=0)
print("training score:", train_score)
print("test score: ", test_score)
print("number of features used: ", coeff_used)
lasso001 = Lasso(alpha=0.01, max_iter=10e5)
lasso001.fit(train_features,train_labels)
train_score001=lasso001.score(train_features,train_labels)
test_score001=lasso001.score(test_features,test_labels)
coeff_used001 = np.sum(lasso001.coef_!=0)
print("training score for alpha=0.01:", train_score001)
print("test score for alpha =0.01: ", test_score001)
print("number of features used: for alpha =0.01:", coeff_used001)
lasso0001 = Lasso(alpha=0.0001, max_iter=10e5)
lasso0001.fit(train_features,train_labels)
train_score0001=lasso0001.score(train_features,train_labels)
test_score0001=lasso0001.score(test_features,test_labels)
coeff_used0001 = np.sum(lasso0001.coef_!=0)
print("training score for alpha=0.0001:", train_score0001)
print("test score for alpha =0.0001: ", test_score0001)
print("number of features used: for alpha =0.0001:", coeff_used0001)

```

```

training score: 0.021558759070389066
test score: 0.020205815285361806
number of features used: 10
training score for alpha=0.01: 0.02156062130858072
test score for alpha =0.01: 0.02020252115769039
number of features used: for alpha =0.01: 10
training score for alpha=0.0001: 0.02156062149479465
test score for alpha =0.0001: 0.020202470244255255
number of features used: for alpha =0.0001: 10

```



```
In [38]: from sklearn.linear_model import Ridge
ridgereg = Ridge(alpha=0, normalize=True)
ridgereg.fit(train_features, train_labels)
y_pred = ridgereg.predict(test_features)
from sklearn import metrics
print("R-Square Value",r2_score(test_labels,y_pred))
print ("mean_absolute_error :",metrics.mean_absolute_error(test_labels, y_pred
))
print ("mean_squared_error : ",metrics.mean_squared_error(test_labels, y_pred
))
print ("root_mean_squared_error : ",np.sqrt(metrics.mean_squared_error(test_la
bels, y_pred)))
ridgereg = Ridge(alpha=0.1, normalize=True)
ridgereg.fit(train_features, train_labels)
y_pred = ridgereg.predict(test_features)
print("R-Square Value",r2_score(test_labels,y_pred))
print ("mean_absolute_error :",metrics.mean_absolute_error(test_labels, y_pred
))
print ("mean_squared_error : ",metrics.mean_squared_error(test_labels, y_pred
))
print ("root_mean_squared_error : ",np.sqrt(metrics.mean_squared_error(test_la
bels, y_pred)))
print(ridgereg.coef_)
```

```
R-Square Value 0.02020246972837836
mean_absolute_error : 11950.144186561989
mean_squared_error : 285134727.5654296
root_mean_squared_error : 16885.932830774545
R-Square Value 0.020080346450857656
mean_absolute_error : 11939.191499945779
mean_squared_error : 285170267.1400758
root_mean_squared_error : 16886.985140636436
[-1.83744566e+02  5.12760030e+01 -2.35563745e+03 -1.96128702e+01
 1.30913268e+03  2.40050122e+03  3.91052677e+02 -3.56696578e+01
-2.03058346e-01  8.41517289e+02]
```

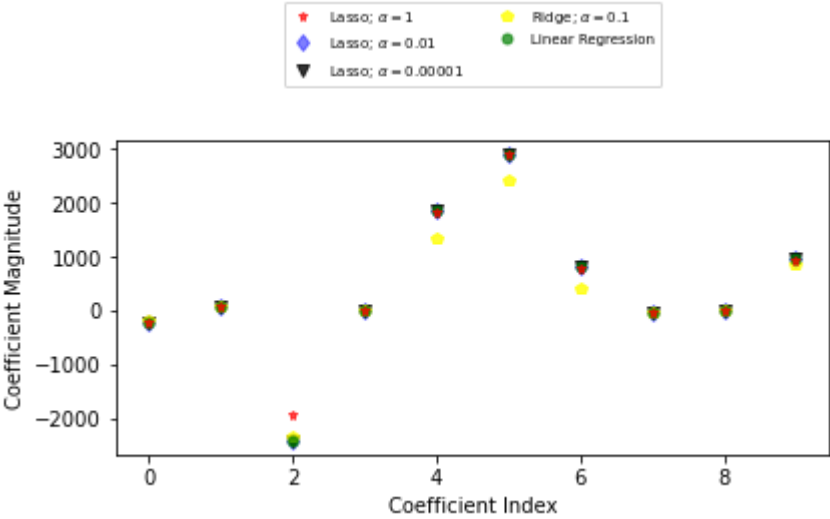


```

In [39]: plt.plot(lasso.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color=
'red',label=r'Lasso;  $\alpha = 1$ ',zorder=7)
plt.plot(lasso001.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color=
'blue',label=r'Lasso;  $\alpha = 0.01$ ')
plt.xlabel('Coefficient Index')
plt.ylabel('Coefficient Magnitude')
plt.legend(fontsize=7,loc="lower center", bbox_to_anchor=(0.5, 1.15), ncol=2)
plt.show()

plt.plot(lasso.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color=
'red',label=r'Lasso;  $\alpha = 1$ ',zorder=7)
plt.plot(lasso001.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color=
'blue',label=r'Lasso;  $\alpha = 0.01$ ')
plt.plot(lasso0001.coef_,alpha=0.8,linestyle='none',marker='v',markersize=6,color=
'black',label=r'Lasso;  $\alpha = 0.0001$ ')
plt.plot(model.coef_,alpha=0.7,linestyle='none',marker='o',markersize=5,color=
'green',label='Linear Regression',zorder=2)
plt.xlabel('Coefficient Index')
plt.ylabel('Coefficient Magnitude')
plt.legend(fontsize=7,loc="lower center", bbox_to_anchor=(0.5, 1.15), ncol=2)
plt.show()

plt.plot(lasso.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color=
'red',label=r'Lasso;  $\alpha = 1$ ',zorder=7)
plt.plot(lasso001.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color=
'blue',label=r'Lasso;  $\alpha = 0.01$ ')
plt.plot(lasso0001.coef_,alpha=0.8,linestyle='none',marker='v',markersize=6,color=
'black',label=r'Lasso;  $\alpha = 0.0001$ ')
plt.plot(ridgereg.coef_,alpha=0.8,linestyle='none',marker='p',markersize=6,color=
'yellow',label=r'Ridge;  $\alpha = 0.1$ ')
plt.plot(model.coef_,alpha=0.7,linestyle='none',marker='o',markersize=5,color=
'green',label='Linear Regression',zorder=2)
plt.xlabel('Coefficient Index')
plt.ylabel('Coefficient Magnitude')
plt.legend(fontsize=7,loc="lower center", bbox_to_anchor=(0.5, 1.15), ncol=2)
plt.tight_layout()
plt.show()

```



# Loaistic Rearession

```
In [34]: from sklearn.linear_model import LogisticRegression
        from sklearn.metrics import classification_report, confusion_matrix
```

```
In [35]: x=result_df.drop('Response', axis = 1)
        y=result_df['Response']
```

```
In [36]: from sklearn.model_selection import train_test_split
        train_features, test_features, train_labels, test_labels = train_test_split(x,
        y, test_size = 0.3, random_state = 0)
```

```
In [37]: from imblearn.under_sampling import NearMiss
        nr = NearMiss()
        train_features, train_labels = nr.fit_resample(train_features, train_labels)
```

```
In [38]: print('Training Features Shape:', train_features.shape)
        print('Training Labels Shape:', train_labels.shape)
        print('Testing Features Shape:', test_features.shape)
        print('Testing Labels Shape:', test_labels.shape)
```

```
Training Features Shape: (65378, 11)
Training Labels Shape: (65378,)
Testing Features Shape: (152444, 11)
Testing Labels Shape: (152444,)
```

```
In [39]: from sklearn.linear_model import LogisticRegression

        logreg = LogisticRegression(max_iter=1000)

        logreg.fit(train_features, train_labels)

        y_pred=logreg.predict(test_features)

        print(classification_report(test_labels, y_pred))
```

	precision	recall	f1-score	support
0	0.94	0.30	0.46	138423
1	0.11	0.82	0.19	14021
accuracy			0.35	152444
macro avg	0.53	0.56	0.32	152444
weighted avg	0.87	0.35	0.43	152444

```
In [40]: from sklearn import metrics
        print("Accuracy:",metrics.accuracy_score(test_labels, y_pred)*100)
```

```
Accuracy: 35.086982760882684
```

```
In [41]: print("Accuracy:", metrics.accuracy_score(test_labels, y_pred))
print("Precision:", metrics.precision_score(test_labels, y_pred))
print("Recall:", metrics.recall_score(test_labels, y_pred))
```

```
Accuracy: 0.35086982760882685
Precision: 0.1067787664700599
Recall: 0.8224805648669853
```

```
In [42]: from sklearn import metrics
cnf_matrix = metrics.confusion_matrix(test_labels, y_pred)
cnf_matrix
```

```
Out[42]: array([[41956, 96467],
               [ 2489, 11532]], dtype=int64)
```

```
In [43]: import numpy as np
predictions = logreg.predict(test_features)
errors = abs(predictions - test_labels)
print('Mean Absolute Error:', round(np.mean(errors), 2), 'degrees.')
```

```
Mean Absolute Error: 0.65 degrees.
```

```
In [44]: from sklearn.feature_selection import RFE
rfe = RFE(logreg, n_features_to_select=8)
fit = rfe.fit(train_features, train_labels)
print("Num Features: %d" % fit.n_features_)
print("Selected Features: %s" % fit.support_)
print("Feature Ranking: %s" % fit.ranking_)
print("Features:", train_features.columns)
```

```
Num Features: 8
Selected Features: [False True True True True True True True False True]
Feature Ranking: [4 1 1 1 1 1 1 1 3 1 2]
Features: Index(['id', 'Gender', 'Age', 'Driving_License', 'Region_Code',
               'Previously_Insured', 'Vehicle_Age', 'Vehicle_Damage', 'Annual_Premium',
               'Policy_Sales_Channel', 'Vintage'],
              dtype='object')
```

```
In [45]: logreg_imp = LogisticRegression(max_iter=1000)
train_important = train_features.drop(['id', 'Annual_Premium', 'Vintage'], axis=1)
test_important = test_features.drop(['id', 'Annual_Premium', 'Vintage'], axis=1)
logreg_imp.fit(train_important, train_labels)
predictions = logreg_imp.predict(test_important)
errors = abs(predictions - test_labels)
print('Mean Absolute Error:', round(np.mean(errors), 2), 'degrees.')
print(classification_report(test_labels, predictions))
```

Mean Absolute Error: 0.38 degrees.

	precision	recall	f1-score	support
0	0.99	0.58	0.73	138423
1	0.19	0.97	0.32	14021
accuracy			0.62	152444
macro avg	0.59	0.77	0.52	152444
weighted avg	0.92	0.62	0.69	152444

```
In [46]: from sklearn import metrics
cnf_matrix = metrics.confusion_matrix(test_labels, predictions)
cnf_matrix
```

```
Out[46]: array([[80271, 58152],
               [ 488, 13533]], dtype=int64)
```


```
In [47]: from sklearn import metrics
print("Accuracy:", metrics.accuracy_score(test_labels, predictions)*100)
```

Accuracy: 61.53341554931647

```
In [48]: print("Accuracy:", metrics.accuracy_score(test_labels, predictions))
print("Precision:", metrics.precision_score(test_labels, predictions))
print("Recall:", metrics.recall_score(test_labels, predictions))
```

Accuracy: 0.6153341554931647  
Precision: 0.1887842644904792  
Recall: 0.9651950645460381

```
In [49]: y_pred_proba = logreg_imp.predict_proba(test_important)[: ,1]
fpr, tpr, _ = metrics.roc_curve(test_labels, y_pred_proba)
auc = metrics.roc_auc_score(test_labels, y_pred_proba)
plt.plot(fpr,tpr,label="data 1, auc="+str(auc))
plt.legend(loc=4)
plt.show()
```



## Assignment 3

## Decision Tree

```
In [56]: import timeit
start = timeit.default_timer()
x=result_df.drop('Response', axis = 1)
y=result_df['Response']
from sklearn.model_selection import train_test_split
train_features, test_features, train_labels, test_labels = train_test_split(x,
y, test_size = 0.3, random_state = 0)
```

```
In [57]: from sklearn.tree import DecisionTreeClassifier
dt = DecisionTreeClassifier()
dt.fit(train_features, train_labels)
dt_pred = dt.predict(test_features)
stop = timeit.default_timer()
print('Time: ', stop - start)
```

Time: 2.2094804000000004

```
In [59]: param_dict = {
    "criterion" : ['gini', 'entropy'],
    "max_depth": range(1,10),
    "min_samples_split": range(1,10),
    "min_samples_leaf": range(1,5)
}
from sklearn.model_selection import GridSearchCV
grid = GridSearchCV(dt,param_grid = param_dict,cv=10,verbose=1,n_jobs=-1)
grid.fit(train_features,train_labels)
```

Fitting 10 folds for each of 648 candidates, totalling 6480 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 52 tasks      | elapsed: 2.6s
[Parallel(n_jobs=-1)]: Done 352 tasks    | elapsed: 18.4s
[Parallel(n_jobs=-1)]: Done 634 tasks    | elapsed: 45.5s
[Parallel(n_jobs=-1)]: Done 984 tasks    | elapsed: 1.6min
[Parallel(n_jobs=-1)]: Done 1434 tasks   | elapsed: 2.8min
[Parallel(n_jobs=-1)]: Done 1984 tasks   | elapsed: 4.6min
[Parallel(n_jobs=-1)]: Done 2634 tasks   | elapsed: 7.5min
[Parallel(n_jobs=-1)]: Done 3384 tasks   | elapsed: 10.8min
[Parallel(n_jobs=-1)]: Done 4234 tasks   | elapsed: 12.4min
[Parallel(n_jobs=-1)]: Done 5184 tasks   | elapsed: 15.6min
[Parallel(n_jobs=-1)]: Done 6234 tasks   | elapsed: 20.8min
[Parallel(n_jobs=-1)]: Done 6480 out of 6480 | elapsed: 22.2min finished
```

```
Out[59]: GridSearchCV(cv=10, estimator=DecisionTreeClassifier(), n_jobs=-1,
    param_grid={'criterion': ['gini', 'entropy'],
    'max_depth': range(1, 10),
    'min_samples_leaf': range(1, 5),
    'min_samples_split': range(1, 10)},
    verbose=1)
```

```
In [60]: grid.best_params_
```

```
Out[60]: {'criterion': 'entropy',
    'max_depth': 8,
    'min_samples_leaf': 3,
    'min_samples_split': 7}
```

```
In [61]: start = timeit.default_timer()
dt = DecisionTreeClassifier(criterion='entropy', max_depth=8, min_samples_leaf
=3, min_samples_split=7)
dt.fit(train_features, train_labels)
dt_pred = dt.predict(test_features)
stop = timeit.default_timer()
print('Time: ', stop - start)
```

Time: 1.1883704000001671



```
In [62]: errors = abs(dt_pred - test_labels)
print('Mean Absolute Error:', round(np.mean(errors), 2), 'degrees.')
print("Accuracy:", metrics.accuracy_score(test_labels, dt_pred)*100)
conf = confusion_matrix(test_labels, dt_pred)
print(conf)
```

Mean Absolute Error: 0.09 degrees.

Accuracy: 90.79596441972134

```
[[138412    11]
 [ 14020     1]]
```

```
In [63]: print("Accuracy:", metrics.accuracy_score(test_labels, dt_pred))
print("Precision:", metrics.precision_score(test_labels, dt_pred))
print("Recall:", metrics.recall_score(test_labels, dt_pred))
```

Accuracy: 0.9079596441972134

Precision: 0.08333333333333333

Recall: 7.132158904500393e-05

```
In [64]: importances = list(dt.feature_importances_)
feature_importances = [(feature, round(importance, 2)) for feature, importance
in zip(x.columns, importances)]
feature_importances = sorted(feature_importances, key = lambda x: x[1], reverse = True)
[print('Variable: {:20} Importance: {}'.format(*pair)) for pair in feature_importances];
```

Variable: Previously_Insured	Importance: 0.55
Variable: id	Importance: 0.28
Variable: Age	Importance: 0.07
Variable: Vehicle_Damage	Importance: 0.07
Variable: Policy_Sales_Channel	Importance: 0.02
Variable: Gender	Importance: 0.0
Variable: Driving_License	Importance: 0.0
Variable: Region_Code	Importance: 0.0
Variable: Vehicle_Age	Importance: 0.0
Variable: Annual_Premium	Importance: 0.0
Variable: Vintage	Importance: 0.0

```
In [66]: rf_most_important = DecisionTreeClassifier(criterion='entropy', max_depth=8, min_samples_leaf=3, min_samples_split=7)
train_important = train_features.loc[:, ['Annual_Premium', 'Vintage', 'Age', 'Region_Code', 'Vehicle_Damage', 'Policy_Sales_Channel', 'Gender', 'Previously_Insured']]
test_important = test_features.loc[:, ['Annual_Premium', 'Vintage', 'Age', 'Region_Code', 'Vehicle_Damage', 'Policy_Sales_Channel', 'Gender', 'Previously_Insured']]
rf_most_important.fit(train_important, train_labels)
predictions = rf_most_important.predict(test_important)
errors = abs(predictions - test_labels)
print('Mean Absolute Error:', round(np.mean(errors), 2), 'degrees.')
```


Mean Absolute Error: 0.09 degrees.

```
In [67]: print("Accuracy:",metrics.accuracy_score(test_labels, predictions))
print("Precision:",metrics.precision_score(test_labels, predictions))
print("Recall:",metrics.recall_score(test_labels, predictions))
```

Accuracy: 0.9079137256959933  
Precision: 0.13043478260869565  
Recall: 0.00021396476713501176

```
In [68]: y_pred_proba = dt.predict_proba(test_features)[::,1]
fpr, tpr, _ = metrics.roc_curve(test_labels, y_pred_proba)
auc = metrics.roc_auc_score(test_labels, y_pred_proba)
print(auc)
plt.plot(fpr,tpr,label="data 1, auc="+str(auc))
plt.legend(loc=4)
plt.title("ROC_AUC Curve")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.show()
```

0.8859370741341136



## Random Forest

```
In [50]: import timeit
start = timeit.default_timer()
x=result_df.drop('Response', axis = 1)
y=result_df['Response']
from sklearn.model_selection import train_test_split
train_features, test_features, train_labels, test_labels = train_test_split(x,
y, test_size = 0.3, random_state = 0)
```

```
In [51]: from sklearn.ensemble import RandomForestClassifier
rf = RandomForestClassifier(n_estimators = 1000, random_state=42)
rf.fit(train_features, train_labels)
predictions = rf.predict(test_features)
stop = timeit.default_timer()
print('Time: ', stop - start)
```

Time: 1021.1808909000006

```
In [52]: errors = abs(predictions - test_labels)
print('Mean Absolute Error:', round(np.mean(errors), 2), 'degrees.')
print("Accuracy:",metrics.accuracy_score(test_labels, predictions))
print("Precision:",metrics.precision_score(test_labels, predictions))
print("Recall:",metrics.recall_score(test_labels, predictions))
conf = confusion_matrix(test_labels,predictions)
print(conf)
```

Mean Absolute Error: 0.1 degrees.

Accuracy: 0.9047781480412479

Precision: 0.39983812221772563

Recall: 0.07046572997646387

```
[[136940  1483]
 [ 13033   988]]
```

```
In [53]: importances = list(rf.feature_importances_)
feature_importances = [(feature, round(importance, 2)) for feature, importance
in zip(x.columns, importances)]
feature_importances = sorted(feature_importances, key = lambda x: x[1], reverse = True)
[print('Variable: {:20} Importance: {}'.format(*pair)) for pair in feature_importances];
```

Variable: id	Importance: 0.26
Variable: Vintage	Importance: 0.19
Variable: Annual_Premium	Importance: 0.17
Variable: Age	Importance: 0.12
Variable: Region_Code	Importance: 0.08
Variable: Vehicle_Damage	Importance: 0.06
Variable: Policy_Sales_Channel	Importance: 0.05
Variable: Previously_Insured	Importance: 0.04
Variable: Gender	Importance: 0.01
Variable: Vehicle_Age	Importance: 0.01
Variable: Driving_License	Importance: 0.0

```
In [54]: print("Accuracy:",metrics.accuracy_score(test_labels, predictions))
print("Precision:",metrics.precision_score(test_labels, predictions))
print("Recall:",metrics.recall_score(test_labels, predictions))
```


Accuracy: 0.9047781480412479

Precision: 0.39983812221772563

Recall: 0.07046572997646387

```
In [55]: y_pred_proba = rf.predict_proba(test_features)[::,1]
fpr, tpr, _ = metrics.roc_curve(test_labels, y_pred_proba)
auc = metrics.roc_auc_score(test_labels, y_pred_proba)
print(auc)
plt.plot(fpr,tpr,label="data 1, auc="+str(auc))
plt.legend(loc=4)
plt.title("ROC_AUC Curve")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.show()
```

0.8871391801108104



## GBM

```
In [87]: import timeit
start = timeit.default_timer()
x=result_df.drop('Response', axis = 1)
y=result_df['Response']
from sklearn.model_selection import train_test_split
train_features, test_features, train_labels, test_labels = train_test_split(x,
y, test_size = 0.3, random_state = 0)
from imblearn.over_sampling import SMOTE
nr = SMOTE()
train_features, train_labels = nr.fit_sample(train_features, train_labels)
```

```
In [88]: from sklearn.ensemble import GradientBoostingClassifier
gbm = GradientBoostingClassifier()
gbm.fit(train_features,train_labels)
predictions = gbm.predict(test_features)
stop = timeit.default_timer()
print('Time: ', stop - start)
```

Time: 97.52382639999996

```
In [ ]: param ={
        "learning_rate" : [1, 0.5, 0.25, 0.1, 0.05, 0.01],
        "n_estimators" : [100, 200]}

grid_search = GridSearchCV(estimator = gbm, param_grid = param,
                           cv = 3, n_jobs = -1, verbose = 2)
grid_search.fit(train_features, train_labels)
grid_search.best_params_
```


```
In [90]: start = timeit.default_timer()
gbm = GradientBoostingClassifier(learning_rate=0.5, n_estimators=200)
gbm.fit(train_features,train_labels)
predictions = gbm.predict(test_features)
stop = timeit.default_timer()
print('Time: ', stop - start)
```

```
In [91]: errors = abs(predictions - test_labels)
print('Mean Absolute Error:', round(np.mean(errors), 2), 'degrees.')
print("Accuracy:",metrics.accuracy_score(test_labels, predictions))
print("Precision:",metrics.precision_score(test_labels, predictions))
print("Recall:",metrics.recall_score(test_labels, predictions))
conf = confusion_matrix(test_labels,predictions)
print(conf)
```

```
Mean Absolute Error: 0.15 degrees.
Accuracy: 0.8491577234919052
Precision: 0.31757196292080014
Recall: 0.5570929320305257
[[121638 16785]
 [ 6210  7811]]
```

```
In [92]: y_pred_proba = gbm.predict_proba(test_features)[::,1]
fpr, tpr, _ = metrics.roc_curve(test_labels, y_pred_proba)
auc = metrics.roc_auc_score(test_labels, y_pred_proba)
print(auc)
plt.plot(fpr,tpr,label="data 1, auc="+str(auc))
plt.legend(loc=4)
plt.title("ROC_AUC Curve")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.show()
```

0.8787520166969816



## XGM

```
In [75]: import timeit
start = timeit.default_timer()
x=result_df.drop('Response', axis = 1)
y=result_df['Response']
from sklearn.model_selection import train_test_split
train_features, test_features, train_labels, test_labels = train_test_split(x,
y, test_size = 0.3, random_state = 0)
from imblearn.over_sampling import SMOTE
nr = SMOTE()
train_features, train_labels = nr.fit_sample(train_features, train_labels)
```

```
In [84]: from xgboost import XGBClassifier as xgb
model_xgb = xgb()
model_xgb.fit(train_features,train_labels)
best_preds = model_xgb.predict(test_features)
stop = timeit.default_timer()
print('Time: ', stop - start)
```

Time: 3162.954396

```
In [ ]: from sklearn.model_selection import GridSearchCV
parameters = {
    "eta"      : [0.05, 0.10, 0.15, 0.20, 0.25, 0.30] ,
    "max_depth" : [ 3, 4, 5, 6, 8, 10, 12, 15]}

grid = GridSearchCV(model_xgb,
                    parameters, n_jobs=-1,
                    scoring="neg_log_loss",
                    cv=3, verbose=2)
grid.fit(train_features, train_labels)
grid.best_params_
```


```
In [ ]: start = timeit.default_timer()
model_xgb = xgb(eta=0.05, max_depth=12)
model_xgb.fit(train_features, train_labels)
best_preds = model_xgb.predict(test_features)
stop = timeit.default_timer()
print('Time: ', stop - start)
```

```
In [85]: errors = abs(best_preds - test_labels)
print('Mean Absolute Error:', round(np.mean(errors), 2), 'degrees.')
print("Accuracy:", metrics.accuracy_score(test_labels, best_preds))
print("Precision:", metrics.precision_score(test_labels, best_preds))
print("Recall:", metrics.recall_score(test_labels, best_preds))
conf = confusion_matrix(test_labels, best_preds)
print(conf)
```

```
Mean Absolute Error: 0.15 degrees.
Accuracy: 0.8471897877253286
Precision: 0.3198943525207799
Recall: 0.5874046073746523
[[120913  17510]
 [  5785   8236]]
```

```
In [86]: y_pred_proba = model_xgb.predict_proba(test_features)[::,1]
fpr, tpr, _ = metrics.roc_curve(test_labels, y_pred_proba)
auc = metrics.roc_auc_score(test_labels, y_pred_proba)
print(auc)
plt.plot(fpr,tpr,label="data 1, auc="+str(auc))
plt.legend(loc=4)
plt.title("ROC_AUC Curve")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.show()
```

0.8813100371610658



## Neural Network Classifier

```
In [1]: from keras import Sequential
from keras.layers import Dense
```

```
In [28]: x=result_df.drop('Response', axis = 1)
y=result_df['Response']
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
x = sc.fit_transform(x)
```

```
In [29]: from sklearn.model_selection import train_test_split
from imblearn.over_sampling import SMOTE
train_features, test_features, train_labels, test_labels = train_test_split(x,
y, test_size = 0.3)
nr = SMOTE()
train_features, train_labels = nr.fit_sample(train_features, train_labels)
```



```
In [12]: print('Training Features Shape:', train_features.shape)
print('Training Labels Shape:', train_labels.shape)
print('Testing Features Shape:', test_features.shape)
print('Testing Labels Shape:', test_labels.shape)
```

```
Training Features Shape: (645800, 11)
Training Labels Shape: (645800,)
Testing Features Shape: (152444, 11)
Testing Labels Shape: (152444,)
```

```
In [13]: def build_model():
        classifier = Sequential()
        classifier.add(Dense(6, activation='relu', kernel_initializer='random_normal', input_dim=11))
        classifier.add(Dense(6, activation='relu', kernel_initializer='random_normal'))
        classifier.add(Dense(1, activation='sigmoid', kernel_initializer='random_normal'))
        classifier.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
        return classifier
```

```
In [14]: keras_model = build_model()  
keras_model.fit(train_features,train_labels, batch_size=64, epochs=100)
```

```
Epoch 1/100
10091/10091 [=====] - 16s 2ms/step - loss: 0.3795 -
accuracy: 0.8377
Epoch 2/100
10091/10091 [=====] - 16s 2ms/step - loss: 0.3604 -
accuracy: 0.8446
Epoch 3/100
10091/10091 [=====] - 16s 2ms/step - loss: 0.3580 -
accuracy: 0.8448
Epoch 4/100
10091/10091 [=====] - 16s 2ms/step - loss: 0.3566 -
accuracy: 0.8451
Epoch 5/100
10091/10091 [=====] - 16s 2ms/step - loss: 0.3554 -
accuracy: 0.8454
Epoch 6/100
10091/10091 [=====] - 16s 2ms/step - loss: 0.3547 -
accuracy: 0.8454
Epoch 7/100
10091/10091 [=====] - 16s 2ms/step - loss: 0.3543 -
accuracy: 0.8453
Epoch 8/100
10091/10091 [=====] - 16s 2ms/step - loss: 0.3540 -
accuracy: 0.8453
Epoch 9/100
10091/10091 [=====] - 16s 2ms/step - loss: 0.3540 -
accuracy: 0.8453
Epoch 10/100
10091/10091 [=====] - 16s 2ms/step - loss: 0.3539 -
accuracy: 0.8454
Epoch 11/100
10091/10091 [=====] - 16s 2ms/step - loss: 0.3538 -
accuracy: 0.8455
Epoch 12/100
10091/10091 [=====] - 16s 2ms/step - loss: 0.3537 -
accuracy: 0.8454
Epoch 13/100
10091/10091 [=====] - 16s 2ms/step - loss: 0.3537 -
accuracy: 0.8453
Epoch 14/100
10091/10091 [=====] - 16s 2ms/step - loss: 0.3536 -
accuracy: 0.8454
Epoch 15/100
10091/10091 [=====] - 16s 2ms/step - loss: 0.3535 -
accuracy: 0.8455
Epoch 16/100
10091/10091 [=====] - 16s 2ms/step - loss: 0.3534 -
accuracy: 0.8457
Epoch 17/100
10091/10091 [=====] - 16s 2ms/step - loss: 0.3530 -
accuracy: 0.8456
Epoch 18/100
10091/10091 [=====] - 16s 2ms/step - loss: 0.3528 -
accuracy: 0.8458
Epoch 19/100
10091/10091 [=====] - 16s 2ms/step - loss: 0.3527 -
accuracy: 0.8459
```

```
Epoch 20/100
10091/10091 [=====] - 16s 2ms/step - loss: 0.3526 -
accuracy: 0.8459
Epoch 21/100
10091/10091 [=====] - 16s 2ms/step - loss: 0.3527 -
accuracy: 0.8460
Epoch 22/100
10091/10091 [=====] - 16s 2ms/step - loss: 0.3525 -
accuracy: 0.8460
Epoch 23/100
10091/10091 [=====] - 16s 2ms/step - loss: 0.3524 -
accuracy: 0.8460
Epoch 24/100
10091/10091 [=====] - 16s 2ms/step - loss: 0.3523 -
accuracy: 0.8462
Epoch 25/100
10091/10091 [=====] - 17s 2ms/step - loss: 0.3523 -
accuracy: 0.8461
Epoch 26/100
10091/10091 [=====] - 16s 2ms/step - loss: 0.3523 -
accuracy: 0.8461
Epoch 27/100
10091/10091 [=====] - 16s 2ms/step - loss: 0.3521 -
accuracy: 0.8461
Epoch 28/100
10091/10091 [=====] - 16s 2ms/step - loss: 0.3520 -
accuracy: 0.8462
Epoch 29/100
10091/10091 [=====] - 16s 2ms/step - loss: 0.3519 -
accuracy: 0.8464
Epoch 30/100
10091/10091 [=====] - 16s 2ms/step - loss: 0.3517 -
accuracy: 0.8464
Epoch 31/100
10091/10091 [=====] - 16s 2ms/step - loss: 0.3515 -
accuracy: 0.8467
Epoch 32/100
10091/10091 [=====] - 16s 2ms/step - loss: 0.3515 -
accuracy: 0.8464
Epoch 33/100
10091/10091 [=====] - 16s 2ms/step - loss: 0.3513 -
accuracy: 0.8465
Epoch 34/100
10091/10091 [=====] - 16s 2ms/step - loss: 0.3514 -
accuracy: 0.8466
Epoch 35/100
10091/10091 [=====] - 16s 2ms/step - loss: 0.3511 -
accuracy: 0.8465
Epoch 36/100
10091/10091 [=====] - 16s 2ms/step - loss: 0.3510 -
accuracy: 0.8468
Epoch 37/100
10091/10091 [=====] - 16s 2ms/step - loss: 0.3510 -
accuracy: 0.8467
Epoch 38/100
10091/10091 [=====] - 16s 2ms/step - loss: 0.3508 -
accuracy: 0.8469
```

Epoch 39/100  
10091/10091 [=====] - 17s 2ms/step - loss: 0.3507 -  
accuracy: 0.8472

Epoch 40/100  
10091/10091 [=====] - 17s 2ms/step - loss: 0.3506 -  
accuracy: 0.8471

Epoch 41/100  
10091/10091 [=====] - 16s 2ms/step - loss: 0.3506 -  
accuracy: 0.8471

Epoch 42/100  
10091/10091 [=====] - 16s 2ms/step - loss: 0.3506 -  
accuracy: 0.8471

Epoch 43/100  
10091/10091 [=====] - 16s 2ms/step - loss: 0.3505 -  
accuracy: 0.8470

Epoch 44/100  
10091/10091 [=====] - 16s 2ms/step - loss: 0.3505 -  
accuracy: 0.8472

Epoch 45/100  
10091/10091 [=====] - 17s 2ms/step - loss: 0.3504 -  
accuracy: 0.8471

Epoch 46/100  
10091/10091 [=====] - 16s 2ms/step - loss: 0.3505 -  
accuracy: 0.8471

Epoch 47/100  
10091/10091 [=====] - 16s 2ms/step - loss: 0.3504 -  
accuracy: 0.8471

Epoch 48/100  
10091/10091 [=====] - 16s 2ms/step - loss: 0.3504 -  
accuracy: 0.8473

Epoch 49/100  
10091/10091 [=====] - 16s 2ms/step - loss: 0.3504 -  
accuracy: 0.8472

Epoch 50/100  
10091/10091 [=====] - 17s 2ms/step - loss: 0.3504 -  
accuracy: 0.8471

Epoch 51/100  
10091/10091 [=====] - 16s 2ms/step - loss: 0.3504 -  
accuracy: 0.8472

Epoch 52/100  
10091/10091 [=====] - 16s 2ms/step - loss: 0.3503 -  
accuracy: 0.8472

Epoch 53/100  
10091/10091 [=====] - 16s 2ms/step - loss: 0.3504 -  
accuracy: 0.8472

Epoch 54/100  
10091/10091 [=====] - 16s 2ms/step - loss: 0.3503 -  
accuracy: 0.8472

Epoch 55/100  
10091/10091 [=====] - 17s 2ms/step - loss: 0.3503 -  
accuracy: 0.8472

Epoch 56/100  
10091/10091 [=====] - 16s 2ms/step - loss: 0.3503 -  
accuracy: 0.8473

Epoch 57/100  
10091/10091 [=====] - 16s 2ms/step - loss: 0.3502 -  
accuracy: 0.8474

Epoch 58/100  
10091/10091 [=====] - 16s 2ms/step - loss: 0.3502 - accuracy: 0.8472

Epoch 59/100  
10091/10091 [=====] - 16s 2ms/step - loss: 0.3502 - accuracy: 0.8473

Epoch 60/100  
10091/10091 [=====] - 16s 2ms/step - loss: 0.3503 - accuracy: 0.8471

Epoch 61/100  
10091/10091 [=====] - 17s 2ms/step - loss: 0.3502 - accuracy: 0.8473

Epoch 62/100  
10091/10091 [=====] - ETA: 0s - loss: 0.3502 - accuracy: 0.84 - 17s 2ms/step - loss: 0.3502 - accuracy: 0.8473

Epoch 63/100  
10091/10091 [=====] - 17s 2ms/step - loss: 0.3503 - accuracy: 0.8473

Epoch 64/100  
10091/10091 [=====] - 17s 2ms/step - loss: 0.3502 - accuracy: 0.8474

Epoch 65/100  
10091/10091 [=====] - 17s 2ms/step - loss: 0.3502 - accuracy: 0.8472

Epoch 66/100  
10091/10091 [=====] - 17s 2ms/step - loss: 0.3501 - accuracy: 0.8473

Epoch 67/100  
10091/10091 [=====] - 16s 2ms/step - loss: 0.3502 - accuracy: 0.8473

Epoch 68/100  
10091/10091 [=====] - 17s 2ms/step - loss: 0.3501 - accuracy: 0.8474

Epoch 69/100  
10091/10091 [=====] - 17s 2ms/step - loss: 0.3502 - accuracy: 0.8474

Epoch 70/100  
10091/10091 [=====] - 16s 2ms/step - loss: 0.3502 - accuracy: 0.8473

Epoch 71/100  
10091/10091 [=====] - 16s 2ms/step - loss: 0.3502 - accuracy: 0.8471

Epoch 72/100  
10091/10091 [=====] - 16s 2ms/step - loss: 0.3500 - accuracy: 0.8474

Epoch 73/100  
10091/10091 [=====] - 17s 2ms/step - loss: 0.3501 - accuracy: 0.8473

Epoch 74/100  
10091/10091 [=====] - 17s 2ms/step - loss: 0.3501 - accuracy: 0.8473

Epoch 75/100  
10091/10091 [=====] - 16s 2ms/step - loss: 0.3501 - accuracy: 0.8474

Epoch 76/100  
10091/10091 [=====] - 17s 2ms/step - loss: 0.3501 - accuracy: 0.8472

Epoch 77/100  
10091/10091 [=====] - 17s 2ms/step - loss: 0.3500 -  
accuracy: 0.8473  
Epoch 78/100  
10091/10091 [=====] - 17s 2ms/step - loss: 0.3501 -  
accuracy: 0.8473  
Epoch 79/100  
10091/10091 [=====] - 17s 2ms/step - loss: 0.3501 -  
accuracy: 0.8472  
Epoch 80/100  
10091/10091 [=====] - 17s 2ms/step - loss: 0.3501 -  
accuracy: 0.8474  
Epoch 81/100  
10091/10091 [=====] - 17s 2ms/step - loss: 0.3500 -  
accuracy: 0.8473  
Epoch 82/100  
10091/10091 [=====] - 17s 2ms/step - loss: 0.3500 -  
accuracy: 0.8474  
Epoch 83/100  
10091/10091 [=====] - 17s 2ms/step - loss: 0.3501 -  
accuracy: 0.8474  
Epoch 84/100  
10091/10091 [=====] - 17s 2ms/step - loss: 0.3500 -  
accuracy: 0.8474  
Epoch 85/100  
10091/10091 [=====] - 17s 2ms/step - loss: 0.3500 -  
accuracy: 0.8472  
Epoch 86/100  
10091/10091 [=====] - 17s 2ms/step - loss: 0.3500 -  
accuracy: 0.8474  
Epoch 87/100  
10091/10091 [=====] - 17s 2ms/step - loss: 0.3499 -  
accuracy: 0.8475  
Epoch 88/100  
10091/10091 [=====] - 17s 2ms/step - loss: 0.3499 -  
accuracy: 0.8472  
Epoch 89/100  
10091/10091 [=====] - 17s 2ms/step - loss: 0.3499 -  
accuracy: 0.8474  
Epoch 90/100  
10091/10091 [=====] - 17s 2ms/step - loss: 0.3499 -  
accuracy: 0.8474  
Epoch 91/100  
10091/10091 [=====] - 17s 2ms/step - loss: 0.3499 -  
accuracy: 0.8474  
Epoch 92/100  
10091/10091 [=====] - 17s 2ms/step - loss: 0.3499 -  
accuracy: 0.8473  
Epoch 93/100  
10091/10091 [=====] - 17s 2ms/step - loss: 0.3499 -  
accuracy: 0.8475  
Epoch 94/100  
10091/10091 [=====] - 17s 2ms/step - loss: 0.3499 -  
accuracy: 0.8473  
Epoch 95/100  
10091/10091 [=====] - 17s 2ms/step - loss: 0.3499 -  
accuracy: 0.8474

```

Epoch 96/100
10091/10091 [=====] - 17s 2ms/step - loss: 0.3499 -
accuracy: 0.8472
Epoch 97/100
10091/10091 [=====] - 17s 2ms/step - loss: 0.3499 -
accuracy: 0.8473
Epoch 98/100
10091/10091 [=====] - 17s 2ms/step - loss: 0.3499 -
accuracy: 0.8472
Epoch 99/100
10091/10091 [=====] - 17s 2ms/step - loss: 0.3499 -
accuracy: 0.8473
Epoch 100/100
10091/10091 [=====] - 17s 2ms/step - loss: 0.3499 -
accuracy: 0.8474

```

Out[14]: <tensorflow.python.keras.callbacks.History at 0x122c3422670>

```

In [30]: eval_model=keras_model.evaluate(train_features, train_labels)
eval_model

```

```

20176/20176 [=====] - 40s 2ms/step - loss: 0.3513 -
accuracy: 0.8465

```

Out[30]: [0.35126250982284546, 0.8464527726173401]

```

In [31]: y_pred=keras_model.predict(test_features)
y_pred =(y_pred>0.5)

```

```

In [32]: print("Accuracy:",metrics.accuracy_score(test_labels, y_pred))
print("Precision:",metrics.precision_score(test_labels, y_pred))
print("Recall:",metrics.recall_score(test_labels, y_pred))
cm = confusion_matrix(test_labels, y_pred)
print(cm)

```

```

Accuracy: 0.7590918632415838
Precision: 0.26553637610827857
Recall: 0.9380787037037037
[[102751  35869]
 [   856 12968]]

```



```
In [33]: from keras.wrappers.scikit_learn import KerasClassifier
from sklearn.metrics import auc
from sklearn.metrics import roc_curve
y_pred_keras = keras_model.predict(test_features).ravel()
fpr_keras, tpr_keras, thresholds_keras = roc_curve(test_labels, y_pred_keras)
auc_keras = auc(fpr_keras, tpr_keras)
plt.plot(fpr_keras, tpr_keras, label="data 1, auc="+str(auc_keras))
plt.legend(loc=4)
plt.title("ROC_AUC Curve")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.show()
```

